

Model Response Evaluation Instructions

Overview

You will evaluate three model-generated research reports that answer a specific research question. Your task is to assess each report's quality across multiple dimensions and provide a final ranking.

Use your assigned username (e.g., [ChattySquirrel-1](#)) to log in. The system will load your assigned question and three model reports.

We have assigned you questions you either picked yourself or have specified as being within your expertise. Your task is to read and grade the model reports as a professor might grade their student's reports.

The evaluation process contains three parts (each of these is described in more detail later):

1. **Constrained Evaluation**

The constrained evaluation is trying to approximate the evaluation process we are automatically conducting over the model results. Your annotation would then serve as a validation of our methods.

2. **Generalized Evaluation**

The generalized assessment serves to get your natural judgement as an expert in the topic on the quality and acceptability of the model answers. The goal of this portion of assessment is to investigate and understand what aspects of writing you value.

3. **Comparative Evaluation**

A brief comparative assessment, where you rank model responses.

4. **Rubric Evaluation (not always available)**

With the evaluation of the model responses fresh in mind, tell us which set of rubrics would be better (useful and comprehensive) for evaluation purposes.

Part 1: Constrained Evaluation

STEP 1. Paragraph-Level Relevance Rating

This task is designed to be tackled as you are reading through the text. As you encounter (U)nnecessary or (O)ff the Mark passages, select the radio box that best suits your judgement.

Background and Introduction	
Label Meaning: (R)elevant, (U)necessary but Acceptable, (O)ff the Mark, or Misleading	
☐ R ☒ U ☐ O	Planning represents a fundamental reasoning challenge in artificial intelligence, involving the state into a desired goal state (Valmeekam et al., 2022). Traditionally studied as an inference systematic reasoning about cause-and-effect relationships and the consequences of actions.
☐ R ☐ U ☒ O	Large Language Models have emerged as powerful tools that can potentially address planning capabilities. These neural networks, containing hundreds of billions of parameters, are trained on vast text corpora (Singh et al., 2023). LLMs learn to model the probability of token generalization when provided with relevant prompts (Singh et al., 2023).
☐ R ☒ U ☐ O	The application of LLMs to planning tasks represents a growing area of research, particularly must be converted into executable action sequences (Singh et al., 2023). However, this approach natural language instruction text with executable plans, making prompt design particularly some domains, such as generating action sequences for virtual environments (Jansen, 2020) planning capabilities or are primarily retrieving solutions from their training data (Valmeekam et al., 2022).

- **R (Relevant):** Content is directly relevant in answering the question and/or provides valuable information that helps contextualize the answer. *This is the default.*
- **U (Unnecessary but Acceptable):** Content is tangentially related or not too helpful for answering the question (i.e. adds to the bulk without adding additional information). But unlike the “off the mark” it does not hurt to have the information there.
- **O (Off the Mark or Misleading):** Content is incorrect, misleading, or completely unrelated. These are paragraphs that don't answer or provide wrong answers to the query. The information hurts the answer and shouldn't be there.

Please use your best judgement.

In case this helps, we include examples we give for automatic evaluations. Keep in mind that these are tuned to our evaluator LLMs particularities, so when taking into account what's here, please infer what is most meaningful/useful to you.:

- **Explicitly consider whether something may be indirectly relevant.** For example, if the question is about the conditions of horses in South Africa, a paragraph about general animal welfare in South Africa is potentially relevant while not being precisely about horses. On the other hand, a paragraph about pig welfare in South Africa is probably irrelevant.
- **Note that subtle differences can make the text irrelevant to the query.** For instance, text about scientific survey paper generation is not relevant to a query about automatic paper review generation. Even though they seem related, they are about very different tasks.
- **Useful background in general is relevant.** If the question is about an approach to creating liver-related proteins, some information about liver-related proteins could contextualize other parts of the answer. If a paragraph contextualizes another part of the answer, then it is relevant.

STEP 2. Citation Validity Check

A	B	C	D	E	F	G
username-number	number	model	copy/paste of the claim	How many citations?	How many citations are "supporting"?	Do all citations, together, cover the claim?
ChattySquirrel-2	2	Model 1	Two-stage training paradigms that integrate temporal graph learning with language model fine-tuning have also emerged, particularly for temporal knowledge graph applications, using hybrid graph tokenization to better model temporal patterns within LLMs (Chang et al., 2025) (Abdallah et al., 2025) (Wu et al., 2024)	3	3	partially covers

Definition: A "claim" here is specifically defined as a sentence or a sequence of consecutive sentences that end in shared citations. Here is a real example:

For the text:

The distinction between memorization and genuine learning represents a critical concern in supervised fine-tuning research, with profound implications for model reliability and practical deployment. Comprehensive studies examining memorization patterns across various fine-tuning tasks have revealed striking disparities in how different types of tasks promote memorization versus generalization. Tasks such as summarization and dialogue consistently exhibit high levels of memorization, with models showing substantial retention of training examples during inference[7][13]. In contrast, classification tasks, reading comprehension, and machine translation demonstrate significantly lower memorization rates, suggesting that the cognitive demands and information processing requirements of different tasks fundamentally influence how models internalize training data. The relationship between memorization and task complexity appears to follow predictable patterns based on the information density requirements of different applications. Tasks requiring models to understand and retain detailed information from inputs naturally lead to higher memorization rates, as the models must encode more comprehensive representations of the training data to perform successfully. Research has identified this phenomenon through analysis of attention patterns, where tasks with high memorization requirements show dense attention distributions across input tokens, while low-memorization tasks demonstrate more selective attention patterns[13].

The following could be considered "claims":

- "Comprehensive studies examining memorization patterns across various fine-tuning tasks have revealed striking disparities in how different types of tasks promote memorization versus generalization. Tasks such as summarization and dialogue consistently exhibit high levels of memorization, with models showing substantial retention of training examples during inference[7][13]." **This would be a claim with 2 citations.**
- "Research has identified this phenomenon through analysis of attention patterns, where tasks with high memorization requirements show dense attention distributions across input tokens, while low-memorization tasks demonstrate more selective attention patterns[13]." **This would be a claim**

with 1 citation.

- Depending on your reading, you may end up picking the whole second portion of the paragraph as your claim. This is acceptable too (as there are no intervening citations). Here you are assuming that the final citation is the bit that captures all of the claims in the text: "In contrast, classification tasks, reading comprehension, and machine translation demonstrate significantly lower memorization rates, suggesting that the cognitive demands and information processing requirements of different tasks fundamentally influence how models internalize training data. The relationship between memorization and task complexity appears to follow predictable patterns based on the information density requirements of different applications. Tasks requiring models to understand and retain detailed information from inputs naturally lead to higher memorization rates, as the models must encode more comprehensive representations of the training data to perform successfully. Research has identified this phenomenon through analysis of attention patterns, where tasks with high memorization requirements show dense attention distributions across input tokens, while low-memorization tasks demonstrate more selective attention patterns[13]." **This would be a claim with 1 citation**

Note that, if you start seeing citation total counts of over, say, 3-5, you may be capturing a “claim” that is longer than how we currently define it.

In this task you will:

- Select 5 claims from the report. We encourage you to (1) select claims that are central to the report (i.e., essential for answering the question) and (2) select claims from a few different sections.
 - When selecting claims, you could also prioritize those with citations that appear with frequency in the report
 - If the report is very short, then stopping at 3 claims is acceptable.
 - Please do not select claims supported by “LLM Memory”. “LLM Memory” means that the model did not offer specific citation (i.e. claim supported by the model’s parametric memory).
- Fill out the Google Spreadsheet you received from us. Specifically, fill out all of the following fields starting from Column B (see example on row 2):
 - Enter in the **number** of the particular evaluation (username-[#])
 - Column A will populate with your username-number automatically
 - The **model** you are evaluating (Model 1, 2, or 3)
 - Copy/Paste a **claim**.
 - Annotations:
 - **How many citations?**
Simply count up the number of citations you see appear with the claim.
 - **How many of the citations are “supporting”?**

- We specifically define “supporting” to mean that the citation supplied is appropriate for the claim, and supplies support to the claim in some way (full or partial portions of the claim).
- Relying ONLY on your prior knowledge -or- what information the tool gives you (e.g., titles, snippets, author list), for each citation decide if it supports (or is likely to support, if you aren’t sure) the claim made.
- How to handle evaluation when snippet/extract is missing:
 - If the citation only gives you a title (rather than a snippet) and you do not have prior knowledge of the paper that it cites, then base your annotation on whether it is *likely* a paper with the title would cover the topic.

When judging with title only, then award a value of 0.5 credit instead of a full 1 if supporting.

- When a citation lacks snippets AND title, then this citation is an automatic 0. For example
- Example:

“Text text text (author1, year1) (author2, year2) (author3, year3)”

 - (author1, year1) has a snippet and checks out. That’s a full 1 credit
 - (author2, year2) has no hover (or the hover over is entirely blank). This is a 0 credit.
 - (author3, year3) this has only title and you have no prior knowledge of this paper, then .5 if likely to support.
 - Total 1.5 credit out of possible 3
- You are not required to follow the link. Only use what knowledge you got or is given to you. Make your best expert guess if need be.
- Write in the number of citations you believe are “supporting”.
- If among the claims you evaluate you see a citation that has no snippet, no title or nothing that will clue you in as to its contents, then this would be considered non supporting.
- **Do all citations (together) fully cover the claim?**
 - Now you are looking at the question from the vantage of the claim coverage.
 - Considering all of the citations together, do they fully cover the claim? Your choices are “fully cover”, “partially cover”, or “does not cover”.
- Optional: Use the **Notes** column if you feel there’s something you need to communicate to us about the particular claim you are looking at.
- **A walkthrough.** Let’s say you come across the following passage:

Robotics and Embodied AI

- **Embodied Agent Navigation:** LLM-Planner enables embodied agents to follow natural language instructions in visually-perceived environments, building on hierarchical planning approaches used in complex vision-and-language navigation tasks like ALFRED (Song et al., 2022) (Shridhar et al., 2019) (Blukis et al., 2021) (Min et al., 2021)

- You choose to pull the first claim in this section, and paste the following into the Column D in the spreadsheet:
“LLM-Planner enables embodied agents to follow natural language instructions in visually-perceived environments, building on hierarchical planning approaches used in complex vision-and-language navigation tasks like ALFRED (Song et al., 2022) (Shridhar et al., 2019) (Blukis et al., 2021) (Min et al., 2021)”
- You note that there are 4 citations associated, and type it into Column E.
- You read the snippets associated with each of the citations:
 - You tally up your supporting counts as you read the snippets, titles or whatever is supplied to you. If using titles only, and you don't know the specific contents, judge based on how /likely/ the citation is to support.
 - Song et al: Very much on target: it is about visual navigation and the use of hierarchical planning approaches to overcome complexities. I consider this “supporting”
 - Shridhar et al: This paper is on ALFRED. It is a vision language navigation benchmark for testing natural instruction navigation.. So this definitely seems to be “supporting”
 - Blukis et al: Uses hierarchical reasoning and tests on ALFRED. “supporting”.
 - For the sake of having a negative example, i'm just going to say Min et al was “not supporting”
 - that's 3 (of 4) that's supporting. You write 3 into Column F.
- Now stepping back, this claim suggests
 - LLM Planning enables embodied agents to follow nl instructions
 - it is about vision language (visual environments)
 - specifically focused on hierarchical planning approaches
 - over complex visual navigation like the challenge posed by ALFRED
 - To me it looks like 3 citations provided touches on all aspects of the claim. We'll call this “fully cover”

- In the Evaluations section, select “Done” for the first question.

Citation Validity Check

Please use the Google sheet to fill out your validity checks. Once you have done this, click Done, and continue

☒ Select ☐ Done

Part 2: Generalized Evaluation

For this part, we ask you to grade the report along several dimensions.

These dimensions were pulled from debrief sessions from prior runs. Thank you to those annotators who have contributed to identifying and defining what makes a good report “good”.

Generalized Evaluation

CITATION GENERAL: Based on your expert knowledge, did all citations you expected to see in an answer appear in the response?

☒ Select ☐ Yes ☐ Partially Yes ☐ No, not what I expected

ANSWER: Did you find an answer in the response?

☒ Select ☐ Yes ☐ Partially Yes ☐ No

CONTENT: Did the report cover all aspects of the answer you expected for it to cover?

☒ Select ☐ Yes ☐ Partially Yes ☐ No, not what I expected

INFORMATIVENESS: Given what it DID cover, how informative was the report?

☒ Select ☐ A (superb) ☐ B (great) ☐ C (acceptable) ☐ D (lacking) ☐ F (fail)

BACKGROUND: How useful was the it offered background for understanding the central answer?

☒ Select ☐ A (superb) ☐ B (great) ☐ C (acceptable) ☐ D (lacking) ☐ F (fail)

ORGANIZATION: How well was the report structured?

☒ Select ☐ A (superb) ☐ B (great) ☐ C (acceptable) ☐ D (lacking) ☐ F (fail)

FLOW: How well did the paragraphs and sections flow into the next?

☒ Select ☐ A (superb) ☐ B (great) ☐ C (acceptable) ☐ D (lacking) ☐ F (fail)

Citation General:

This is a question you have expertise on. This means you have a general sense for what works should have been cited. This question gets at whether or not the report addressed those.

You marking “No, not what I expected” does not mean necessarily commenting on the validity of the citations provided by the report. It means there were works that you would have absolutely used, and it doesn’t appear in the report.

Answer Quality: Sometimes models do provide plentiful information and details, but it doesn’t answer the very question the user is asking. So, this rating is to assess whether the model did actually answer the question.

Content: Again, as an expert, you have a general sense for what an answer should cover for it to be complete. This question assesses the coverage of what you expected.

Informativeness: This assesses the level of detail or depth of information supplied. For example, a report might have hit all the right topic notes, but did not supply enough information on any particular one (informativeness is lacking). Conversely a short well-written report could also prove to be highly informative. For this rating, please focus on the information that is most relevant to the question.

Background: Models always provide background information, but it isn't always useful. This is gauging the utility of the background information in supporting or contextualizing the actual answer.

Organization: This is related to the next question "Flow". This question is directed at the *structure* of the response. Is the response using structure in the best way possible? For example, are the sections meaningful, are the section titles useful, are the more important parts of the answers organizationally highlighted?

Flow: This question is directed at how well a section or a paragraph flows into the next section or paragraph.

For organization and flow, please do not deliberate too heavily to try to tease them apart. For some reports, the distinction will be obvious, for others, not so obvious. So please simply give us your best evaluation.

Part 3: Final Ranking and Debrief

Holistic Ranking

Assign medals to the three reports. Two way ties are acceptable (one gold and two silvers for example). Please resist assigning three way ties. *Holistic Ranking should capture your preference, first and foremost.*

-  **Gold:** Best overall report
-  **Silver:** Second best report
-  **Bronze:** Third best report

Ranking Reasoning

(This is a non optional field. If we find later that it is missing, we will ask you to fill it out.)

Provide a brief explanation of your ranking decisions (1-2 sentences are sufficient). Consider:

- Why did you rank as you did?

- Why did one edge over the other?

Part 4: Rubric Comparison (Not always available)

Once you are done with the final ranking, we ask you to look through **three** sets of rubrics and evaluate the utility of the rubric for assessing the content of the answers you have just looked at.

Note the selectable 3 tabs, which will bring up the rubric and the questions associated with the rubric. Each of the rubric criteria or “ingredients” are sorted by their weight, ordered by the most important (highest weight) to the least important (lowest weight). *Note that the weights are proportional to the total number of ingredients, so the weights cannot be compared across the rubrics.*

You can double click on any of the ingredients to see the full text.

Rubric A Rubric B Rubric C

Rubric A

Sorted by most important (highest weight) to least important (lowest weight).

Weight	Ingredient
0.11	Discuss communication efficiency challenges as a major bottleneck in federated learning. This can include c
0.11	Explain privacy and security concerns in federated learning. This may include risks of adversarial attacks, in
0.11	Address data heterogeneity challenges, particularly non-IID data distribution. This can include differences i
0.11	Discuss system heterogeneity challenges. Examples can include variability in hardware and network capabi
0.05	Discuss scalability challenges. This can include issues arising when managing or aggregating models across
0.05	Address fairness and bias issues in federated learning. Examples include unequal data contributions across
0.05	Discuss resource management challenges. This can include uneven computational and energy constraints,

About ingredients:

For many of the ingredients you will see: a sentence that introduces the criterion and a sentence that specifies the possible examples.

Please note that you are *not* responsible for evaluating the comprehensiveness of the examples given in the ingredients.

For each rubric answer 3 questions.

- **What are its strengths?**
- **Are there any missing criteria?**

As an expert who is familiar with the details of the questions, are there ingredients you’d have wanted to see but are missing?

- **What are its weaknesses?**

Generally,

- A good rubric is one that:
 - Includes “ingredients” (rubric items) that are important to cover
 - Generally is comprehensive and does not leave out critical information
 - Usable as is without too many heavy edits (small edits such as wording is expected)

Aspects of a rubric that you might consider:

- appropriateness of the ingredient for the given question
- utility of ingredient in evaluating the response to the question
- comprehensiveness of the criteria
- appropriate levels of detailedness
- specificity of the ingredients
- (overall take on) weight assignment of the ingredients

Finally assign gold, silver, bronze (ties allowed; but no 3 way ties please) to the rubrics.

Please rank the three rubrics:

Rubric A: Raking

Rubric B: Raking

Rubric C: Raking

Not Yet Rated



Not Yet Rated



Not Yet Rated



On Snippets and Tables:

- Your prior feedback told us the tables were either not very useful or hard to view/read. We have converted the table into lists (The title will say TABLE if it is that). Hopefully this will make things easier to navigate.
- Snippets (gray box) are viewable when hovering over the yellow highlighted citations.

amental reasoning challenge in artificial intelligence, involving the generation of action sequences that transform an initial

te (Valmeekam et al., 2022)

: cause-and-effect relations

re emerged as powerful tool

etworks, containing hundre

ora (Singh et al., 2023) . LLM

ed with relevant prompts (S

planning tasks represents a

cutable action sequences (S

...with some additional ...

"Most claims about LLM planning capabilities are however based on common sense tasks-where it becomes hard to tell whether LLMs are planning or merely retrieving from their vast world knowledge. There is a strong need for systematic and extensible planning benchmarks with sufficient diversity to evaluate whether LLMs have innate planning capabilities [...]Our studies also show that on many critical capabilities-including plan generation-LLM performance falls quite short, even with the SOTA models [...]Planning is a reasoning task that has been well studied in the AI community. In its most basic form, planning involves coming up with a course of actions (policy) which when executed would take an agent from a certain initial state to a desired world state. Planning has generally been studied primarily as an inference problem on world and reward models."

If you see no highlights over a citation, then no snippets or citation details were supplied by the model. When this happens, quite often, you will see the citation in a non-standard form. In the following example you see an unlinked citation in blue. This is an example of

a raw citation that we don't have additional detail for.

Modern approaches to evaluating document organization recognize distinct measurement categories. These include form-based measures, which focus on grammatical correctness, structural integrity, and overall coherence, as well as content-based measures that assess precision and recall in information representation (Ba-Alwi et al. , 2015). The complexity of document organization is further reflected in the multi-dimensional nature of assessment, which considers factors such as topical distribution patterns and reader engagement [53835829 | Amin et al. | 2019 | Citations: 21].

If you see a yellow highlighted citation without a hover text, this is a bug; please let us know.

Technical Notes

- **Save Early and Often:** Your progress is saved automatically for individual reports, but save the debrief before switching windows
- **Reset to Login:** will clear all progress. Save before resetting.
- **Save Buttons:**
 - **“Save Report Evaluations”** (at the end of each of the individual report evals) are interlinked. You click on one, it will save any information you have input for all three model evaluations. It doesn't hurt though to push it multiple times.
 - **“Save Debrief”** and **“Save Rubric Comparison”** are saved separately.

-